





Jun Tian


<https://tianjun.me/about>
Software Engineer @ Microsoft, Beijing


Three Messages

- **Julia is easy to interact with other languages**
- **Julia is productive to provide HTTP services**
- **Julia is FAST and easy to use for Reinforcement Learning!**










Write the engine in pure Julia



Write a wrapper in pure Julia



C/C++

ccall



C++

CxxWrap.jl

TextWorld

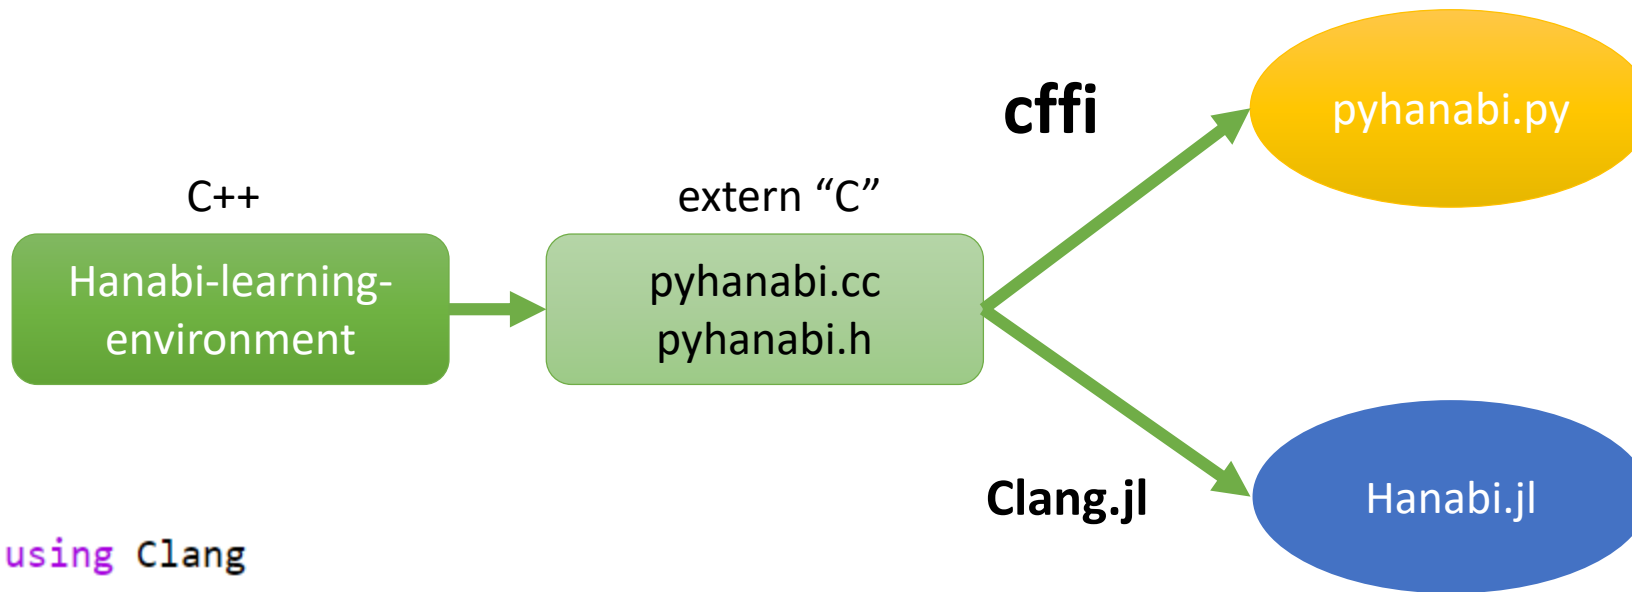
Python

PyCall.jl

.....

.....

.....



using Clang



```
const HANABI_H = joinpath(@__DIR__, "..", "deps", "usr", "include", "pyhanabi.h") |> normpath
```

```
wc = init(; headers = [HANABI_H],  
           output_file = joinpath(@__DIR__, "libhanabi_api.jl"),  
           common_file = joinpath(@__DIR__, "libhanabi_common.jl"),  
           clang_includes = [CLANG_INCLUDE],  
           header_wrapped = (root, current) -> root == current,  
           header_library = x->"libpyhanabi",  
           clang_diagnostics = true)
```


```
run(wc)
```




finalizer



HTTP.WebSockets.WebSocket




```
julia> play()
[system]:A new player [951983957397953247] joined the game!
[system]:A new player [11946767017492306892] joined the game!
[system]:GAME STARTS!
Life tokens: 3
Info tokens: 8
Fireworks: R0 Y0 G0 W0 B0
Hands:
Cur player
XX || XX|RYGWB12345
XX || XX|RYGWB12345
XX || XX|RYGWB12345
XX || XX|RYGWB12345
XX || XX|RYGWB12345
-----
R1 || XX|RYGWB12345
B4 || XX|RYGWB12345
W4 || XX|RYGWB12345
W5 || XX|RYGWB12345
Y3 || XX|RYGWB12345
Deck size: 40
Discards:
[11946767017492306892]:Hello World!
msg >hi
[951983957397953247]:hi
msg >RevealRank(1,1)
[system]:Player 951983957397953247 takes an action [RevealRank(1, 1)]
```



+ Interact.jl




+





How to train a smart agent?

Rainbow

- **Distributional RL(C51)**
- **Prioritized Replay**
- **Multi-step Learning**
- Dueling Networks
- Double Q-learning
- Noise Nets



State



Reward Distribution

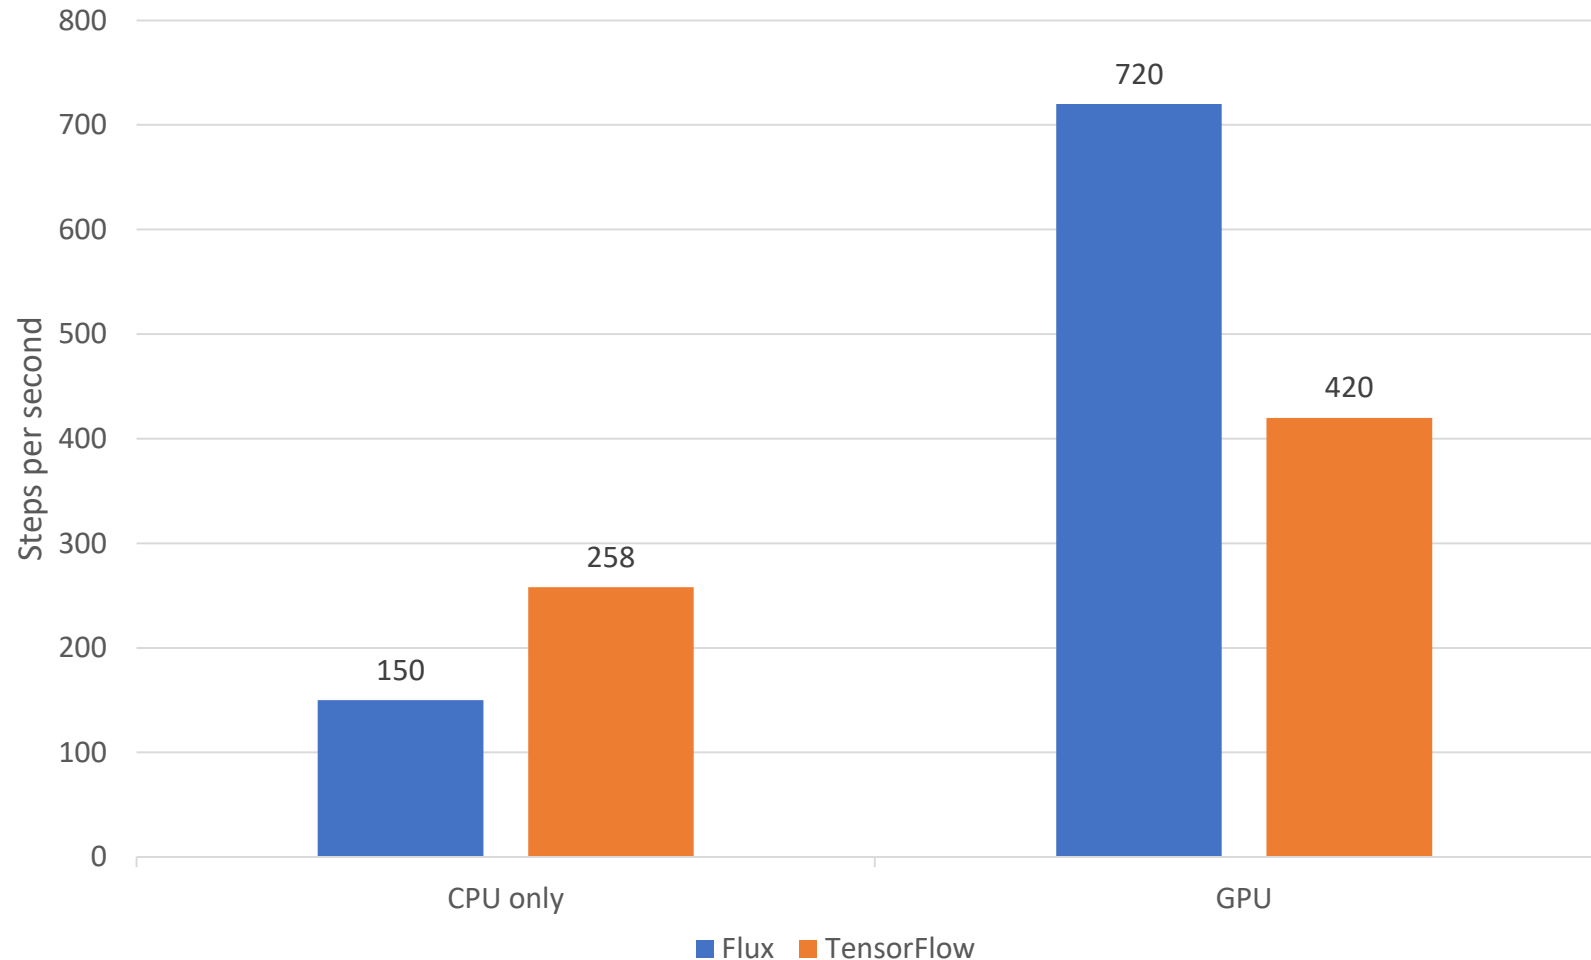


Actions

```
Chain(
    Dense(658, 512),
    Dense(512, 512),
    Dense(512, 20 * 51)
)
```

softmax

Performance Comparison Between Flux and TensorFlow



Environment:

Azure Standard_NC12s_v3

V100 card

Intel(R) Xeon(R) CPU E5-2690 v4 @ 2.60GHz (x12)

Ubuntu 18.04

CUDA 10.0

CUDNN 7.6.0

Julia v1.1.1

Flux v0.8.3


Tracker v0.2.2

Python v3.7.3

TensorFlow v1.13.1

Q1: Faster on GPU but slower on CPU?

Performance Comparison Between Flux and TensorFlow



Environment:

Azure Standard_NC12s_v3

V100 card

Intel(R) Xeon(R) CPU E5-2690 v4 @ 2.60GHz (x12)

Ubuntu 18.04

CUDA 10.0

CUDNN 7.6.0

Julia v1.1.1


Flux v0.8.3

Tracker v0.2.2

Python v3.7.3

TensorFlow v1.13.1

Q1: Faster on GPU but Slower on CPU?



Environment:

Azure Standard_NC12s_v3

V100 card

Intel(R) Xeon(R) CPU E5-2690 v4 @ 2.60GHz (x12)

Ubuntu 18.04

CUDA 10.0

CUDNN 7.6.0

Julia v1.1.1

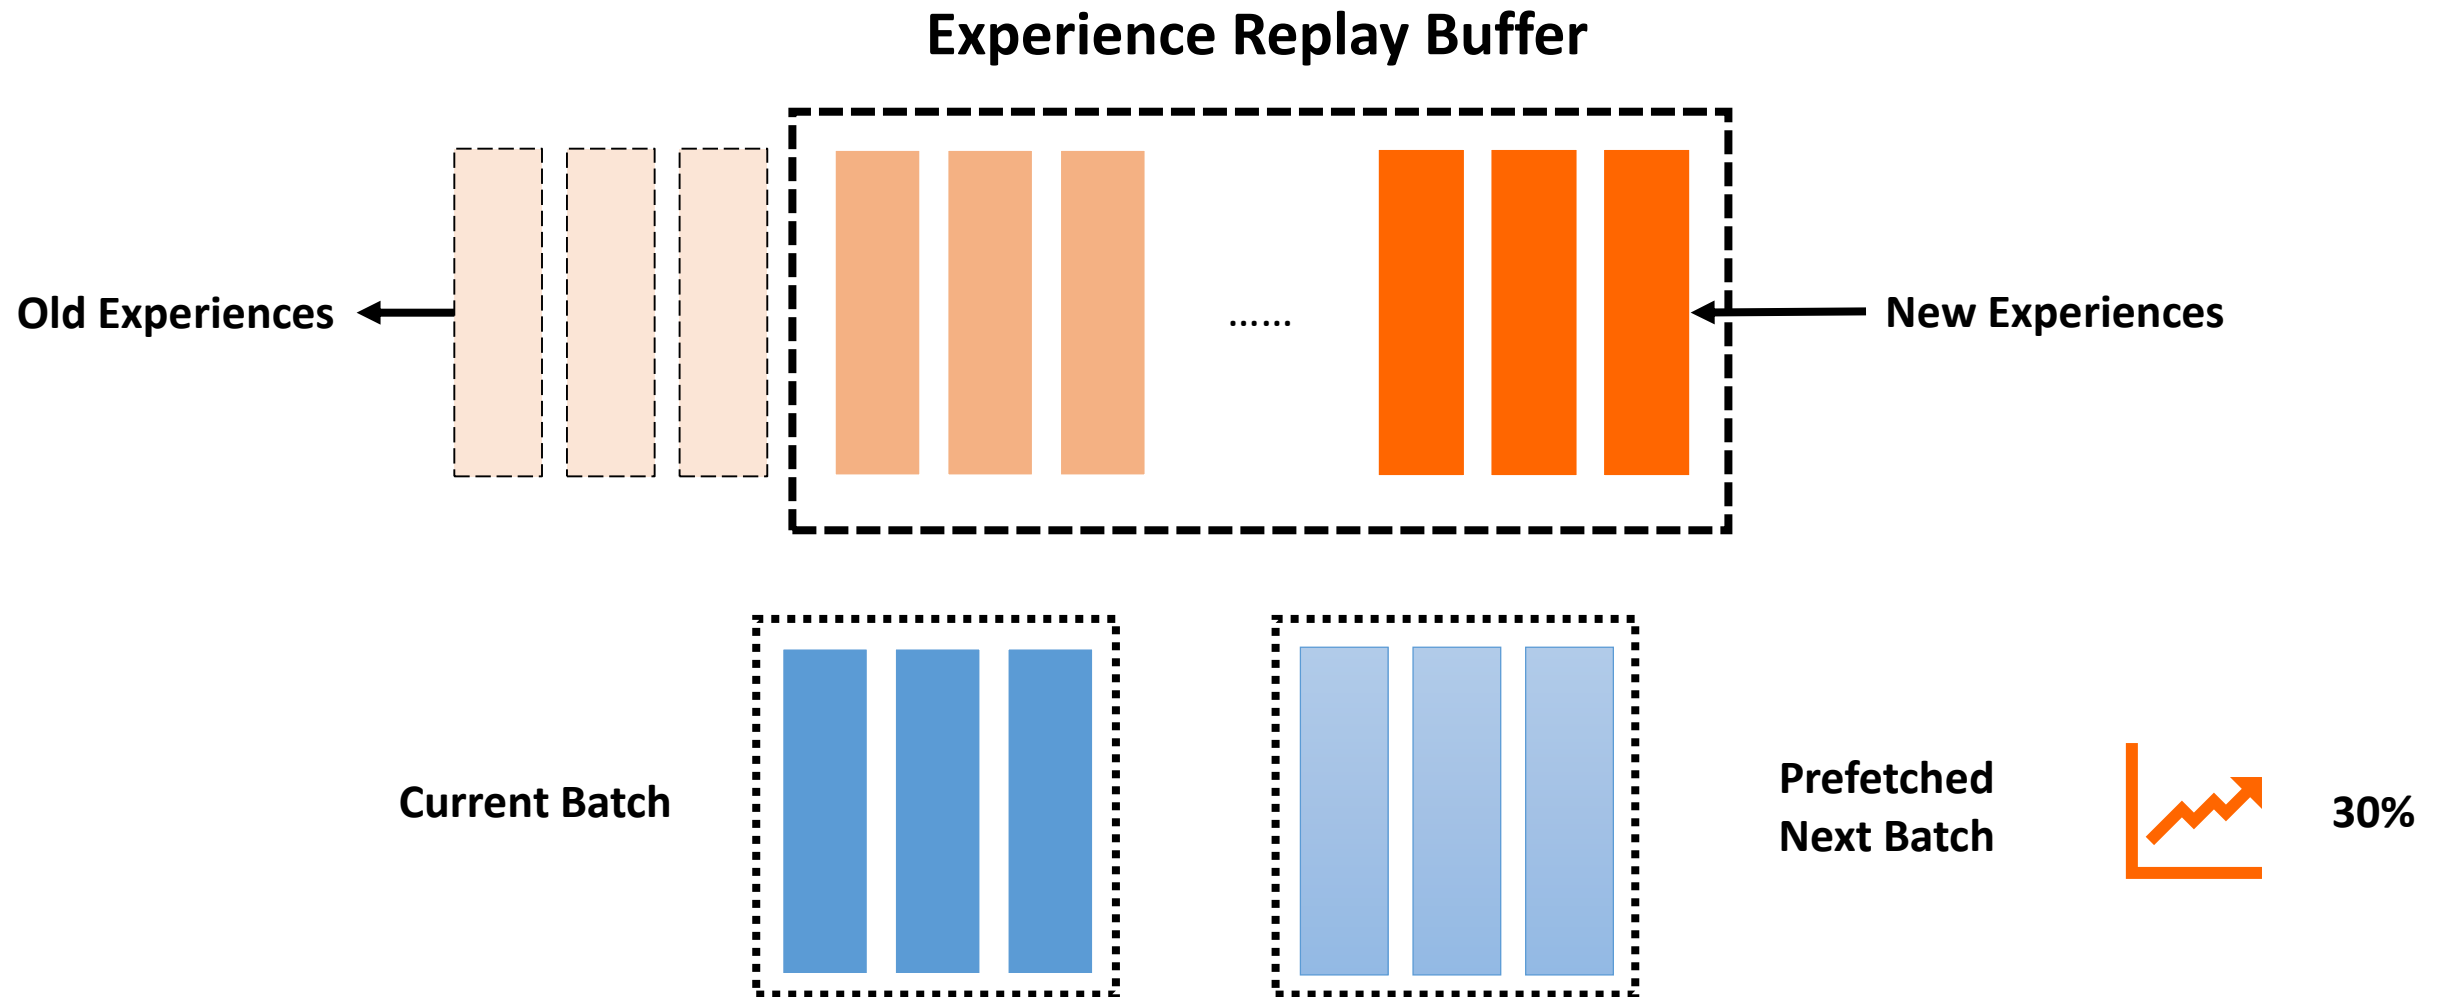
Flux v0.8.3

Tracker v0.2.2

Python v3.7.3


TensorFlow v1.13.1

tf.contrib.staging.StagingArea



Q2: Special Tricks on GPU?

Performance Comparison Between Flux and TensorFlow



Environment:

Azure Standard_NC12s_v3

V100 card

Intel(R) Xeon(R) CPU E5-2690 v4 @ 2.60GHz (x12)

Ubuntu 18.04

CUDA 10.0

CUDNN 7.6.0

Julia v1.1.1

Flux v0.8.3

Tracker v0.2.2

Python v3.7.3





TensorFlow v1.13.1

Tricks:

using CuArrays

It's Just 3 Dense Layers
Except ...

$$P^\pi Z$$



Source: A Distributional Perspective on Reinforcement Learning (Figure 1)

Fused Broadcast

$$(\Phi \hat{\mathcal{T}} Z_{\theta}(x, a))_i = \sum_{j=0}^{N-1} \left[1 - \frac{||[\hat{\mathcal{T}} z_j]_{V_{\text{MIN}}}^{V_{\text{MAX}}} - z_i|}{\Delta z} \right]_0^1 p_j(x', \pi(x'))$$

Python

```
reshaped_target_support = tf.reshape(reshaped_target_support, [batch_size, num_dims, 1])
numerator = tf.abs(tiled_support - reshaped_target_support)
quotient = 1 - (numerator / delta_z)
clipped_quotient = tf.clip_by_value(quotient, 0, 1)
weights = weights[:, None, :]
inner_prod = clipped_quotient * weights
projection = tf.reduce_sum(inner_prod, 3)
```




Julia

```
clamp.(1 .- abs.(tiled_support .- target_support) ./ delta_z, 0, 1) .* weights
```



@views

Experience Replay Buffer



Q3: The Impact of Hardware?

Performance Comparison Between Flux and TensorFlow



Environment:

Azure Standard_NC12s_v3

V100 card

Intel(R) Xeon(R) CPU E5-2690 v4 @ 2.60GHz (x12)

Ubuntu 18.04

CUDA 10.0

CUDNN 7.6.0

Julia v1.1.1


Flux v0.8.3

Tracker v0.2.2

Python v3.7.3

TensorFlow v1.13.1

Performance Comparison Between Flux and TensorFlow



Environment:

RTX 2080 TI card

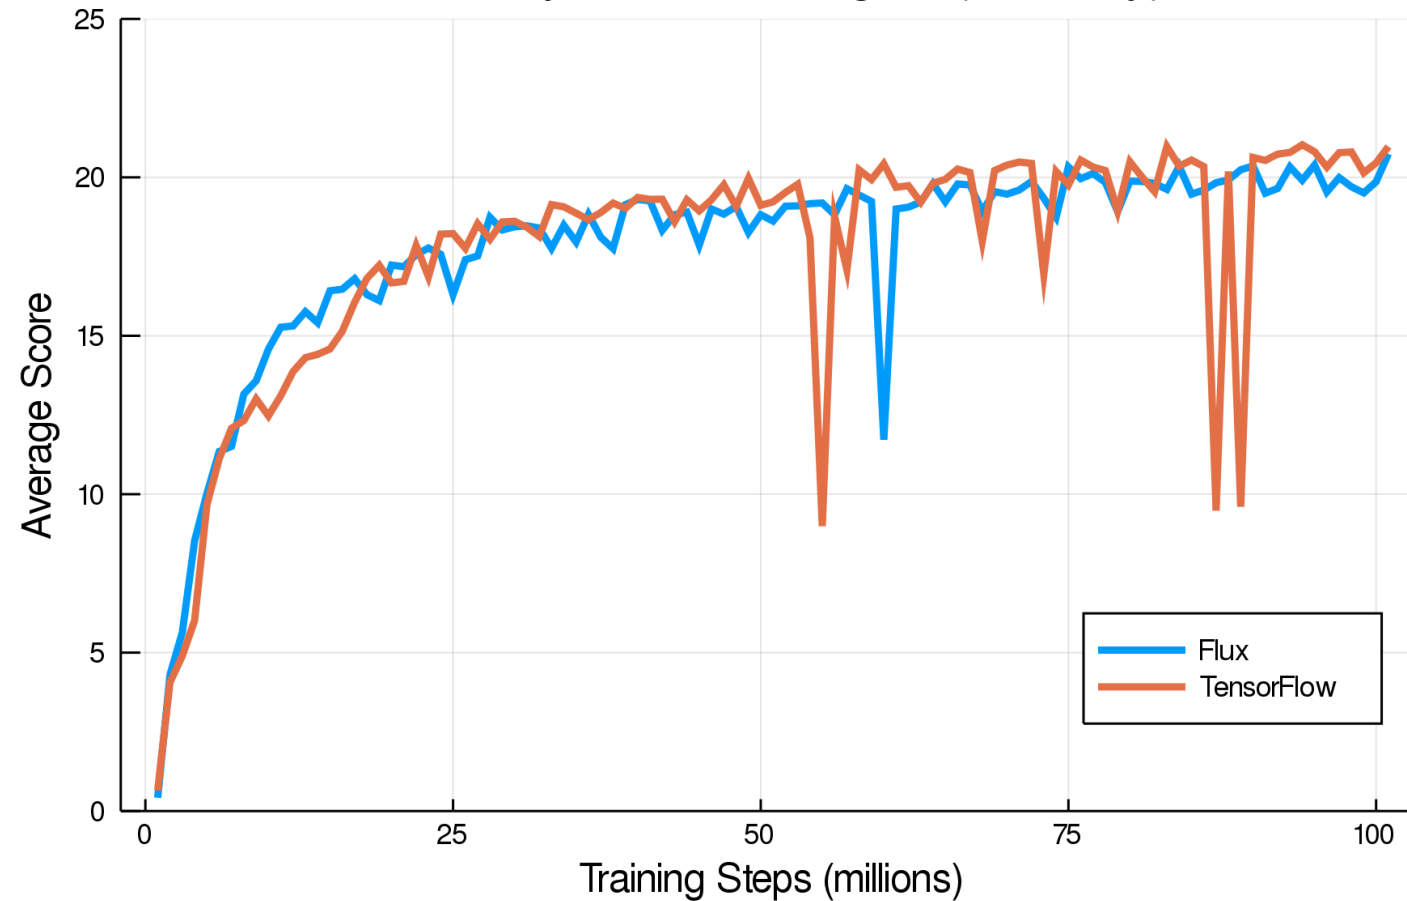
Intel(R) Xeon(R) W-2123 CPU
@ 3.60GHz(×8)

Ubuntu 18.04
CUDA 10.0
CUDNN 7.6.0

Julia v1.1.0
Flux v0.8.3
Tracker v0.2.2

Python v3.7.3
TensorFlow v1.13.1

Two Player Rainbow Agent (Self Play)



Training Time with 2080TI:

Flux: \approx **1** day

TensorFlow: \approx **5** days



“... I was wondering, ... if you are willing to share it (pretrained model) with us. Our computation power is limited, and it would be very expensive for us training many agents...”

How about giving **julia** a try and get some free speedup?



Takeaway Messages

- Remember to release resources with a finalizer
- Avoid extra allocation with @view
- Fused broadcast will make your code fast and easy to read

Reference

For More Details:

- https://tianjun.me/essays/Lets_Play_Hanabi
 - Bayesian Action Decoder
 - Distributed Prioritized Experience Replay

Some Useful Links:

- <https://github.com/JuliaReinforcementLearning/Hanabi.jl>
- <https://github.com/deepmind/hanabi-learning-environment>
- <https://arxiv.org/abs/1710.02298> (Rainbow)
- <https://github.com/JuliaReinforcementLearning/ReinforcementLearning.jl>
- <https://github.com/Ju-jl/ReinforcementLearningAnIntroduction.jl>

More Questions?

Special thanks to

- @Huda
- @Roger
- @Gnimuc
- @jbrea

for their guide and inspiration